



Texture Quilting (Due Saturday 2/23/2019)

In this assignment, you will develop code to stitch together image patches sampled from an input texture in order to synthesize new texture images. You can download the test image used to generate the example above from assignment folder Canvas.

You should start by reading through the whole assignment, looking at the provided code in detail to make sure you understand what it does. The main function ***quilt_demo*** appears at the end. You will need to write several subroutines in order for it to function properly.

Name: *Simrik Manandhar*

SID: *32369856*

1. Shortest Path [25 pts]

Write a function ***shortest_path*** that takes an 2D array of ***costs***, of shape HxW, as input and finds the *shortest vertical path* from top to bottom through the array. A vertical path is specified by a single horizontal location for each row of the H rows. Locations in successive rows should not differ by more than 1 so that at each step the path either goes straight or moves at most one pixel to the left or right. The cost is the sum of the costs of each entry the path traverses. Your function should return an length H vector that contains the index of the path location (values in the range 0..W-1) for each of the H rows.

You should solve the problem by implementing the dynamic programming algorithm described in class. You will have a for-loop over the rows of the "cost-to-go" array (M in the slides), computing the cost of the shortest path up to that row using the recursive formula that depends on the costs-to-go for the previous row. Once you have get to the last row, you can then find the smallest total cost. To find the path which actually has this smallest cost, you will need to do backtracking. The easiest way to do this is to also store the index of whichever minimum was selected at each location. These indices will also be an HxW array. You can then backtrack through these indices, reading out the path.

Finally, you should create at least three test cases by hand where you know the shortest path and see that the code gives the correct answer.

In [1]:

```
#modules used in your code
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
def shortest_path(costs):
    """
    This function takes an array of costs and finds a shortest path from the
    top to the bottom of the array which minimizes the total costs along the
    path. The path should not shift more than 1 location left or right between
    successive rows.

    In other words, the returned path is one that minimizes the total cost:

        total_cost = costs[0,path[0]] + costs[1,path[1]] + costs[2,path[2]] + ..
        .

    subject to the constraint that:

        abs(path[i]-path[i+1])<=1

    Parameters
    -----
    costs : 2D float array of shape HxW
            An array of cost values

    Returns
    -----
    path : 1D array of length H
           indices of a vertical path. path[i] contains the column index of
           the path for each row i.
    """

    rows = costs.shape[0]
    columns = costs.shape[1]
    path=list()
    #creating an array with infinity values and with extra columns for padding
    path_cost= np.full((rows,columns+2),np.inf)
    #appending the costs array in path_cost
    path_cost[:,1:columns+1] = costs

    for i in range(1,rows):
        path_cost[i,1:columns+1] = path_cost[i,1:columns+1] + np.min(np.stack((p
ath_cost[i-1,0:columns], path_cost[i-1,1:columns+1],path_cost[i-1,2:columns+2]))
,axis=0)

    j=np.where(path_cost[rows-1]==min(path_cost[rows-1]))
    path.append(j[0][0]-1)
```

```

        for i in range(rows-2,-1,-1):
            j = np.where(path_cost[i] == min(path_cost[i][j[0][0]],path_cost[i][j[0]
[0]-1],path_cost[i][j[0][0]+1]))
            path.append(j[0][0]-1)
        path.reverse()

    return path

```

In [3]:

```

#
# your test code goes here.  come up with at least 3 test cases
#

costs1 = np.array([[2,3,4,1],[1,4,5,4],[2,3,1,2],[4,3,1,4]])
path1 = shortest_path(costs1)

print(costs1)
print(path1)

costs2 = np.array([[1,2,4,3],[2,1,5,6],[6,4,1,1],[2,5,5,6]])
path2 = shortest_path(costs2)

print(costs2)
print(path2)

costs3 = np.array([[2,3,3,9,1],[3,2,6,9,1],[1,6,9,6,2],[7,1,1,1,2]])
path3 = shortest_path(costs3)

print(costs3)
print(path3)

```

```

[[2 3 4 1]
 [1 4 5 4]
 [2 3 1 2]
 [4 3 1 4]]
[0, 0, 1, 2]
[[1 2 4 3]
 [2 1 5 6]
 [6 4 1 1]
 [2 5 5 6]]
[0, 1, 1, 0]
[[2 3 3 9 1]
 [3 2 6 9 1]
 [1 6 9 6 2]
 [7 1 1 1 2]]
[4, 4, 4, 3]

```

2. Image Stitching: [25 pts]

Write a function ***stitch*** that takes two gray-scale images, ***left_image*** and ***right_image*** and a specified ***overlap*** and returns a new output image by stitching them together along a seam where the two images have very similar brightness values. If the input images are of widths ***w1*** and ***w2*** then your stitched result image returned by the function should be of width ***w1+w2-overlap*** and have the same height as the two input images.

You will want to first extract the overlapping strips from the two input images and then compute a cost array given by the absolute value of their difference. You can then use your ***shortest_path*** function to find the seam along which to stitch the images where they differ the least in brightness. Finally you need to generate the output image by using pixels from the left image on the left side of the seam and from the right image on the right side of the seam. You may find it easiest to code this by first turning the path into an alpha mask for each image and then using the standard equation for compositing.

In [4]:

```
def stitch(left_image, right_image, overlap):
    """
    This function takes a pair of images with a specified overlap and stitches them
    together by finding a minimal cost seam in the overlap region.

    Parameters
    -----
    left_image : 2D float array of shape HxW1
        Left image to stitch

    right_image : 2D float array of shape HxW2
        Right image to stitch

    overlap : int
        Width of the overlap zone between left and right image

    Returns
    -----
    stitched : 2D float array of shape Hx(W1+W2-overlap)
        The resulting stitched image
    """

    # inputs should be the same height
    assert(left_image.shape[0]==right_image.shape[0])

    # your code here
    mask=left_image.shape[1]+right_image.shape[1]-overlap
    stitched=np.zeros((right_image.shape[0],mask))
```

```

assert(stitched.shape[0]==left_image.shape[0])
assert(stitched.shape[1]==(left_image.shape[1]+right_image.shape[1]-overlap)
)

overlapleft=left_image[:,left_image.shape[1]-overlap+1:left_image.shape[1]]
overlapright=right_image[:,1:overlap]

costs = np.abs(overlapleft-overlapright)
seam = shortest_path(costs)

for i in range (1,stitched.shape[0]):
    temp=seam[i]
    for j in range(1,stitched.shape[1]):
        if (j<(left_image.shape[1]-overlap+seam[i])):
            stitched[i,j]=left_image[i,j]
        else:
            stitched[i,j]=right_image[i,temp]
            temp=temp+1

return stitched

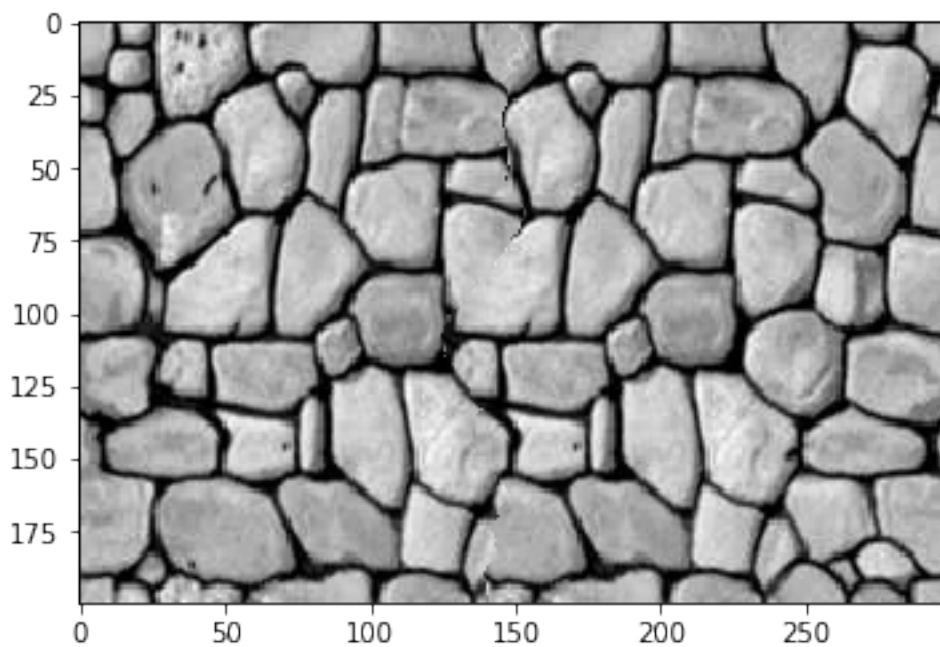
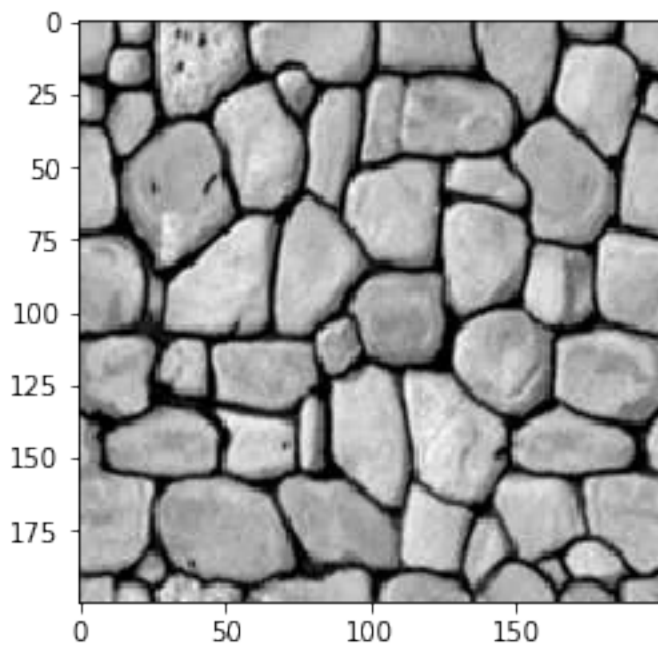
```

In [5]:

```

# Testing stitch function
#load an image
Image = plt.imread('rock_wall.jpg')
#convert to float data type and scale to [0..1] if necessary
if (Image.dtype == np.uint8):
    Image = Image.astype(float) / 256
#if your image is color (shape HxWx3), convert to grayscale by averaging together R,G,B values
Image = np.mean(Image, -1)
#display the image in the notebook using a grayscale colormap
plt.imshow(Image,cmap=plt.cm.gray)
#force matplotlib to go ahead and display the plot now
plt.show()
Image_stitch = stitch(Image,Image,100)
plt.imshow(Image_stitch,cmap=plt.cm.gray)
plt.show()

```



3. Texture Quilting: [25 pts]

Write a function ***synth_quilt*** that takes as input an array indicating the set of texture tiles to use, an array containing the set of available texture tiles, the ***tilesize*** and ***overlap*** parameters and synthesizes the output texture by stitching together the tiles. ***synth_quilt*** should utilize your stitch function repeatedly. First, for each horizontal row of tiles, construct the stitched row by repeatedly stitching the next tile in the row on to the right side of your row image. Once you have row images for all the rows, you can stitch them together to get the final image. Since your stitch function only works for vertical seams, you will want to transpose the rows, stitch them together, and then transpose the result. You may find it useful to look at the provided code below which simply puts down the tiles with the specified overlap but doesn't do stitching. Your quilting function will return a similar result but with much smoother transitions between the tiles.

In [6]:

```
def synth_quilt(tile_map, tiledb, tilesize, overlap):
```

```
    """
```

```
    This function takes as input an array indicating the set of texture tiles
    to use at each location, an array containing the database of available textu
```

re

y

tiles, tileSize and overlap parameters, and synthesizes the output texture by stitching together the tiles

Parameters

tile_map : 2D array of int

Array storing the indices of which tiles to paste down at each output location

tiledb : 2D array of size ntiles x npixels

Collection of sample tiles to select from

tilesize : (int,int)

Size of a tile in pixels

overlap : int

Amount of overlap between tiles

Returns

output : 2D float array

The resulting synthesized texture of size

"""

determine output size based on overlap and tile size

*outh = (tilesize[0]-overlap)*tile_map.shape[0] + overlap*

*outw = (tilesize[1]-overlap)*tile_map.shape[1] + overlap*

output = np.zeros((outh,outw))

The code below is a dummy implementation that pastes down each

tile in the correct position in the output image. You need to

replace this with your own version that stitches each row and then

stitches together the columns

for i **in** range(tile_map.shape[0]):

for j **in** range(tile_map.shape[1]):

 tile_vec = tiledb[tile_map[i,j],:];

 tile_image = np.reshape(tile_vec,tilesize)

if j==0:

 row_image=tile_image

else:

 row_image=stitch(row_image,tile_image,overlap)

if i==0:

 output=row_image.T

else:

 output=stitch(output,row_image.T,overlap)

return output.T

4. Texture Synthesis Demo [25pts]

The function provided below ***quilt_demo*** puts together the pieces. It takes a sample texture image and a specified output size and uses the functions you've implemented previously to synthesize a new texture sample.

You should write some additional code in the cells that follow to in order demonstrate the final result and experiment with the algorithm parameters in order to produce a compelling visual result and write explanations of what you discovered.

Test your code on the provided image *rock_wall.jpg*. There are three parameters of the algorithm. The *tilesize*, *overlap* and *K*. In the provided ***texture_demo*** code below, these have been set at some default values. Include in your demo below images of three example texture outputs when you: (1) increase the tile size, (2) decrease the overlap, and (3) decrease the value for K. For each result explain how it differs from the default setting of the parameters and why.

Test your code on two other texture source images of your choice. You can use images from the web or take a picture of a texture yourself. You may need to resize or crop your input image to make sure that the ***tiledb*** is not overly large. You will also likely need to modify the ***tilesize*** and ***overlap*** parameters depending on your choice of texture. Once you have found good settings for these parameters, synthesize a nice output texture. Make sure you display both the image of the input sample and the output synthesis for your two other example textures in your submitted pdf.

In [7]:

```
#skimage is only needed for sample tiles code provided below
#you should not use it in your own code
import skimage as ski

def sample_tiles(image,tilesize,randomize=True):
    """
        This function generates a library of tiles of a specified size from a given
        source image

        Parameters
        -----
        image : float array of shape HxW
            Input image

        tilesize : (int,int)
            Dimensions of the tiles in pixels

        Returns
        -----
        tiles : float array of shape numtiles x numpixels
            The library of tiles stored as vectors where numpixels is the
            product of the tile height and width
    """
```



```

tiles = ski.util.view_as_windows(image,tilesiz)
ntiles = tiles.shape[0]*tiles.shape[1]
npix = tiles.shape[2]*tiles.shape[3]
assert(npix==tilesiz[0]*tilesiz[1])

print("library has ntiles = ",ntiles,"each with npix = ",npix)

tiles = tiles.reshape((ntiles,npix))

# randomize tile order
if randomize:
    tiles = tiles[np.random.permutation(ntiles),:]

return tiles

def topkmatch(tilestrip,dbstrips,k):
    """
    This function finds the top k candidate matches in dbstrips that
    are most similar to the provided tile strip.

    Parameters
    -----
    tilestrip : 1D float array of length npixels
        Grayscale values of the query strip

    dbstrips : 2D float array of size npixels x numtiles
        Array containing brightness values of numtiles strips in the database
        to match to the npixels brightness values in tilestrip

    k : int
        Number of top candidate matches to sample from

    Returns
    -----
    matches : list of ints of length k
        The indices of the k top matching tiles
    """
    assert(k>0)
    assert(dbstrips.shape[0]>k)
    error = (dbstrips-tilestrip)
    ssd = np.sum(error*error,axis=1)
    ind = np.argsort(ssd)
    matches = ind[0:k]
    return matches

def quilt_demo(sample_image, ntilesout=(10,20), tilesiz=(30,30), overlap=5, k=5
):
    """
    This function takes an image and quilting parameters and synthesizes a
    new texture image by stitching together sampled tiles from the source image.

```

Parameters

sample_image : 2D float array
Grayscale image containing sample texture

ntilesout : list of int
Dimensions of output in tiles, e.g. (3,4)

tilesize : int
Size of the square tile in pixels

overlap : int
Amount of overlap between tiles

k : int
Number of top candidate matches to sample from

Returns

img : list of int of length K
The resulting synthesized texture of size
"""

generate database of tiles from sample
tiledb = sample_tiles(sample_image,tilesize)
number of tiles in the database
nsampletiles = tiledb.shape[0]

if (nsampletiles<k):
 print("Error: tile database is not big enough!")

generate indices of the different tile strips
i,j = np.mgrid[0:tilesize[0],0:tilesize[1]]
top_ind = np.ravel_multi_index(np.where(i<overlap),tilesize)
bottom_ind = np.ravel_multi_index(np.where(i>=tilesize[0]-overlap),tilesize)
left_ind = np.ravel_multi_index(np.where(j<overlap),tilesize)
right_ind = np.ravel_multi_index(np.where(j>=tilesize[1]-overlap),tilesize)

initialize an array to store which tile will be placed
in each location in the output image
tile_map = np.zeros(ntilesout,'int')

print('row:')
for i in range(ntilesout[0]):
 for j in range(ntilesout[1]):

 if (i==0)&(j==0): # first row first tile
 matches = np.zeros(k) #range(nsampletiles)

 elif (i==0): # first row (but not first tile)
 left_tile = tile_map[i,j-1]
 tilestrip = tiledb[left_tile,right_ind]

```
dbstrips = tiledb[:,left_ind]
```

```
matches = topkmatch(tilestrip,dbstrips,k)
```

```
elif (j==0): # first column (but not first row
```

```
above_tile = tile_map[i-1,j]
```

```
tilestrip = tiledb[above_tile,bottom_ind]
```

```
dbstrips = tiledb[:,top_ind]
```

```
matches = topkmatch(tilestrip,dbstrips,k)
```

```
else: # neighbors above and to the left
```

```
left_tile = tile_map[i,j-1]
```

```
tilestrip_1 = tiledb[left_tile,right_ind]
```

```
dbstrips_1 = tiledb[:,left_ind]
```

```
above_tile = tile_map[i-1,j]
```

```
tilestrip_2 = tiledb[above_tile,bottom_ind]
```

```
dbstrips_2 = tiledb[:,top_ind]
```

```
# concatenate the two strips
```

```
tilestrip = np.concatenate((tilestrip_1,tilestrip_2))
```

```
dbstrips = np.concatenate((dbstrips_1,dbstrips_2),axis=1)
```

```
matches = topkmatch(tilestrip,dbstrips,k)
```

```
#choose one of the k matches at random
```

```
tile_map[i,j] = matches[np.random.randint(0,k)]
```

```
output = synth_quilt(tile_map,tiledb,tilesize,overlap)
```

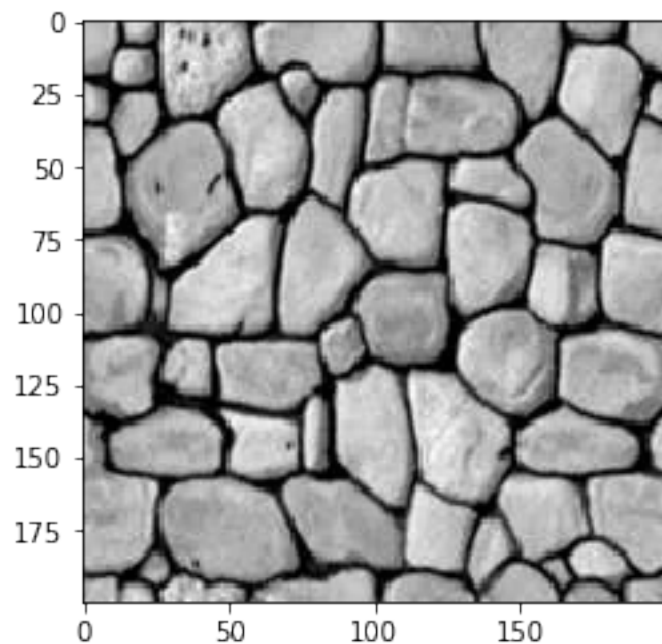
```
return output
```

In [8]:

```
# load in rock_wall.jpg
# run and display results for quilt_demo with
#
# (0) default parameters
# (1) increased tile size
# (2) decrease the overlap
# (3) increase the value for K.

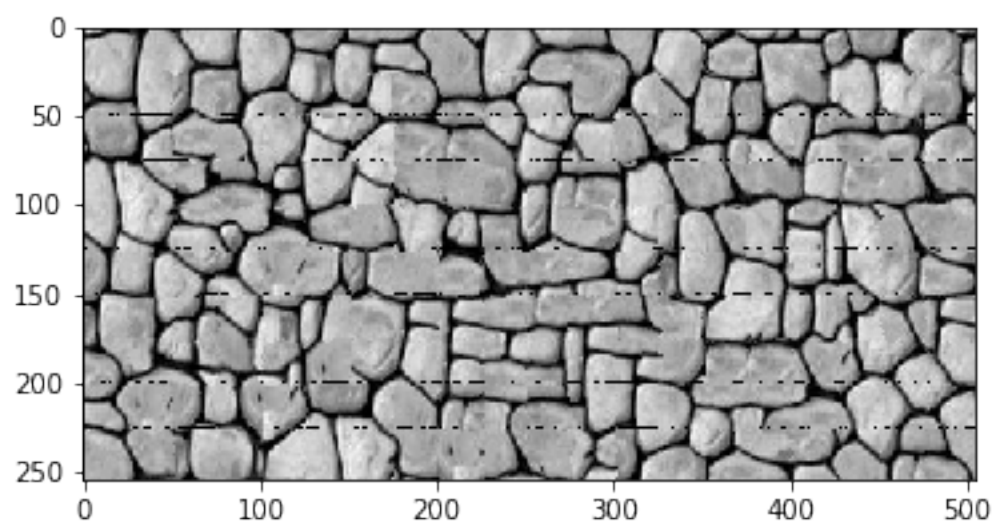
#load an image
I = plt.imread('rock_wall.jpg')
#convert to float data type and scale to [0..1] if necessary
if (I.dtype == np.uint8):
    I = I.astype(float) / 256
#if your image is color (shape HxWx3), convert to grayscale by averaging together R,G,B values
I = np.mean(I, -1)
#display the image in the notebook using a grayscale colormap
plt.imshow(I,cmap=plt.cm.gray)
#force matplotlib to go ahead and display the plot now
print("Original Image:")
plt.show()
# (0) default parameters
print("Default Parameters:")
default = quilt_demo(I)
plt.imshow(default,cmap=plt.cm.gray)
plt.show()
# (1) increased tile size
print("Increased tile size:")
increase_tile = quilt_demo(I,tilesize=(60,50))
plt.imshow(increase_tile,cmap=plt.cm.gray)
plt.show()
# (2) decrease the overlap
print("Decreased overlap:")
decrease_overlap = quilt_demo(I,overlap=3)
plt.imshow(decrease_overlap,cmap=plt.cm.gray)
plt.show()
# (3) increase the value for K.
print("Increased K:")
increase_K = quilt_demo(I,k=8)
plt.imshow(increase_K,cmap=plt.cm.gray)
plt.show()
```

Original Image:



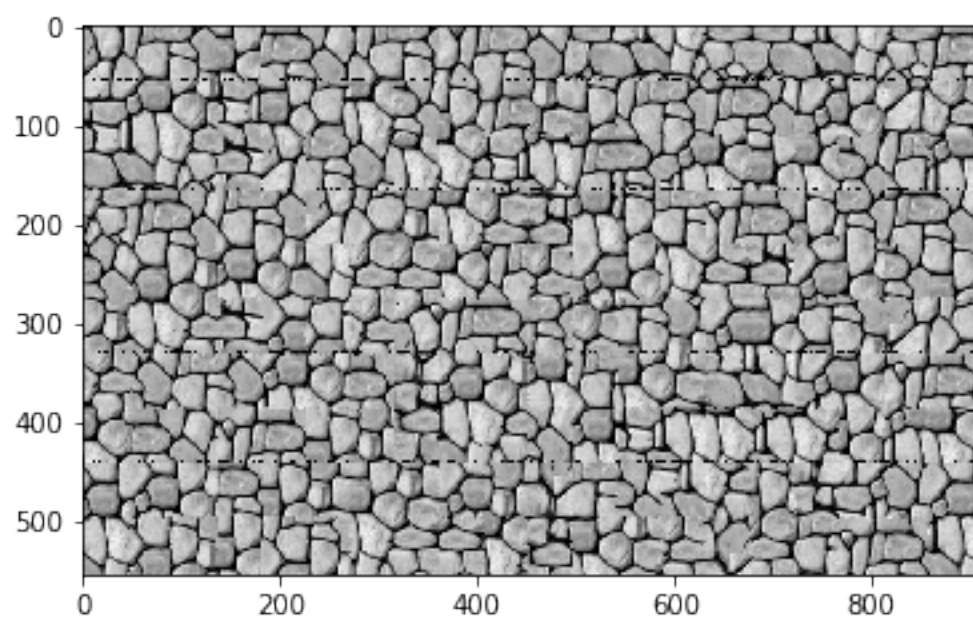
Default Parameters:

library has ntiles = 29241 each with npix = 900
row:



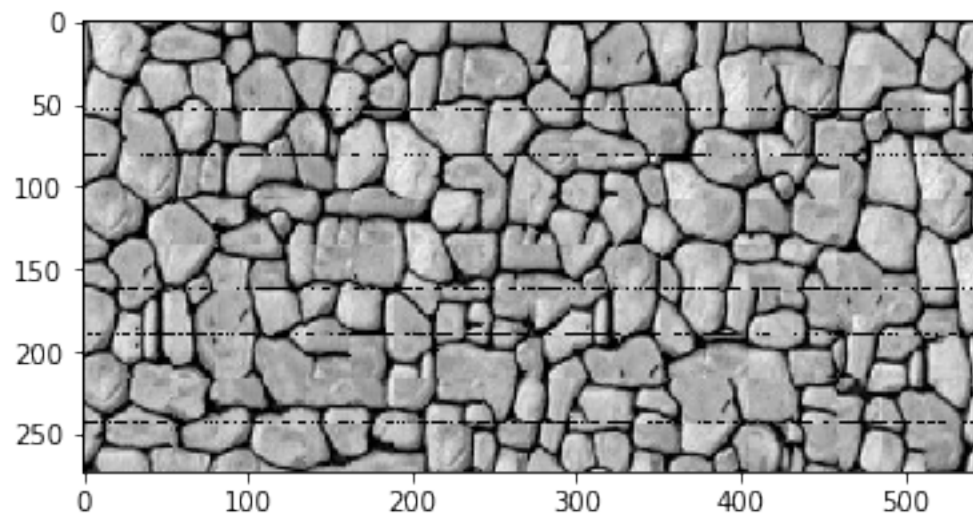
Increased tile size:

library has ntiles = 21291 each with npix = 3000
row:



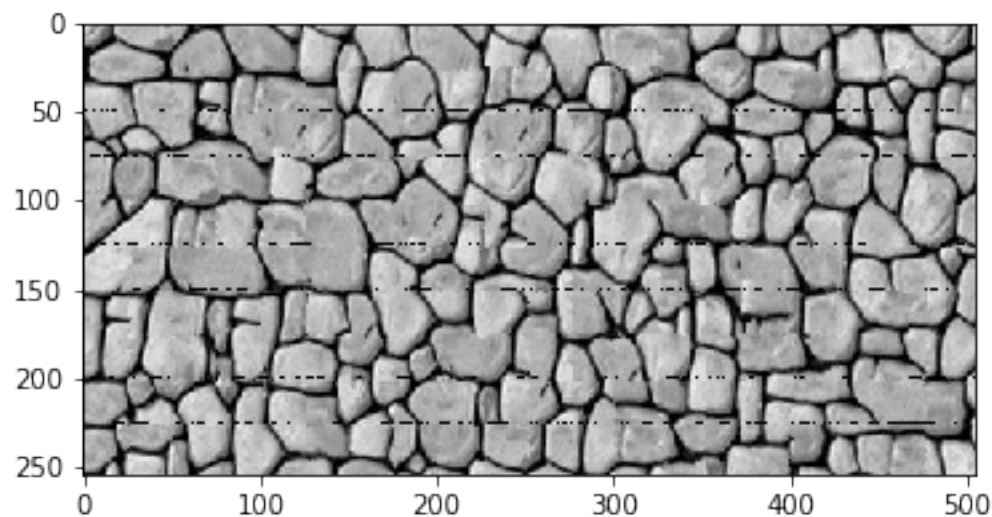
Decreased overlap:

library has ntiles = 29241 each with npix = 900
row:



Increased K:

library has ntiles = 29241 each with npix = 900
row:



For each result shown, explain here how it differs visually from the default setting of the parameters and explain why:

Increasing the tile size: Increasing only the tile size, size of sampling becomes smaller which results in larger number of sampling images with smaller size. However, the resolution of result image is increased but the result image is closely packed stitched image texture.

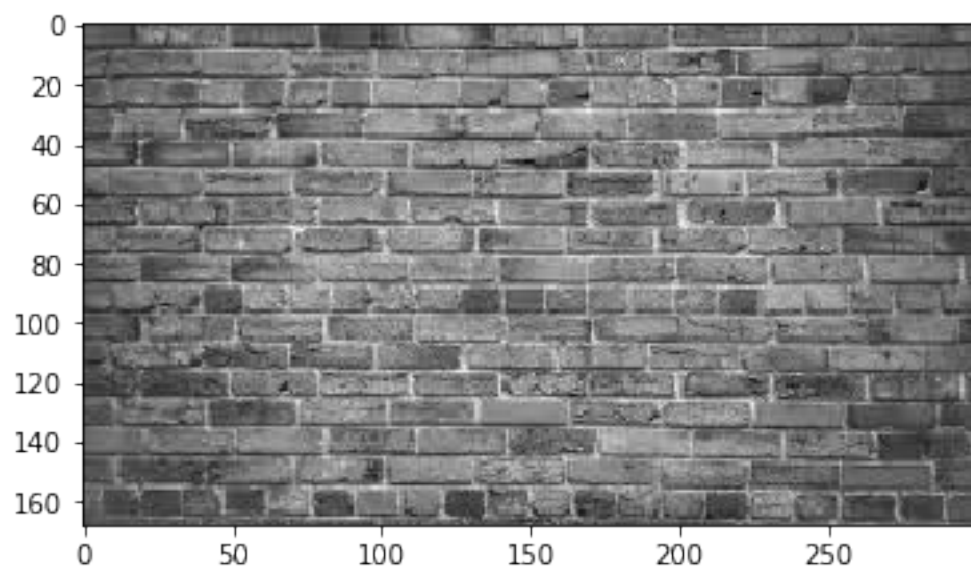
Decreasing overlap: Decreasing overlap between neighbouring tiles somewhat increases the resolution of the final output. The sampling image is also increased. Small overlap number means more detail in the image to work with and visible seam.

Increasing K: K represent the number of top nearest neighbours which matches in the list of tiles. Increasing the value of K means more nearest neighbours can be matched which means there will be more possible way of stitching different tiles together.

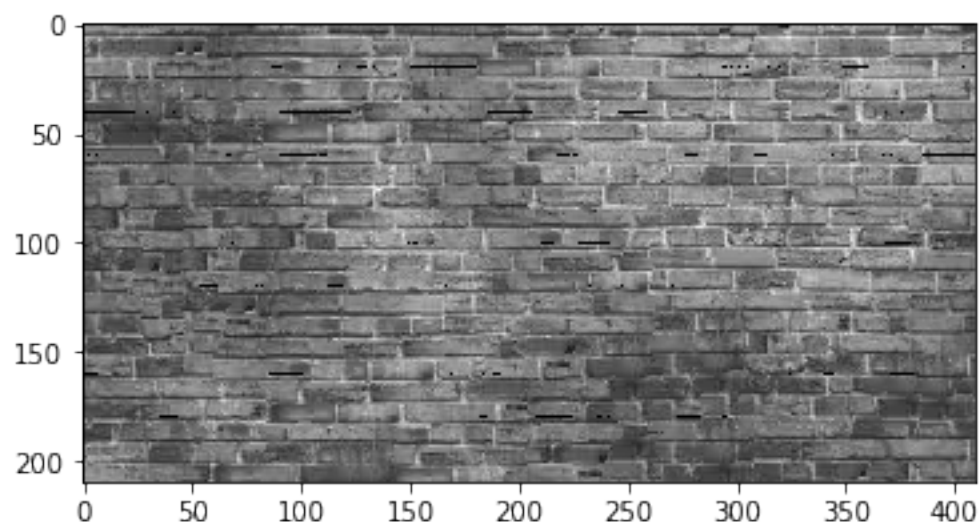
In [9]:

```
# # load in yourimage1.jpg
#
# call quilt_demo, experiment with parameters if needed to get a good result
#
# display your source image and the resulting synthesized texture
#
img1 = plt.imread('img_1.jpeg')
#convert to float data type and scale to [0..1] if necessary
if (img1.dtype == np.uint8):
    img1= img1.astype(float) / 256
img1 = np.mean(img1, -1)
#display the image in the notebook using a grayscale colormap
plt.imshow(img1,cmap=plt.cm.gray)
#force matplotlib to go ahead and display the plot now
print("Original Image:")
plt.show()
img1_stitch = quilt_demo(img1,overlap=10,tilesize=(30,30),k=5)
plt.imshow(img1_stitch,cmap=plt.cm.gray)
plt.show()
```

Original Image:



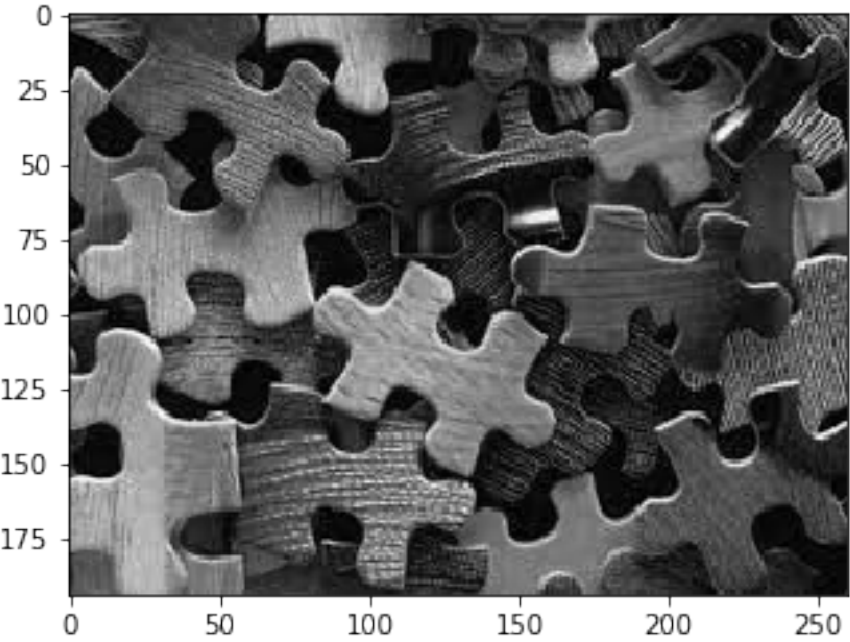
library has ntiles = 37669 each with npix = 900
row:



In [11]:

```
#  
# load in yourimage2.jpg  
#  
# call quilt_demo, experiment with parameters if needed to get a good result  
#  
# display your source image and the resulting synthesized texture  
#  
img2 = plt.imread('img_2.jpeg')  
#convert to float data type and scale to [0..1] if necessary  
if (img2.dtype == np.uint8):  
    img2= img2.astype(float) / 256  
img2 = np.mean(img2, -1)  
#display the image in the notebook using a grayscale colormap  
plt.imshow(img2,cmap=plt.cm.gray)  
#force matplotlib to go ahead and display the plot now  
print("Original Image:")  
plt.show()  
img2_stitch = quilt_demo(img2,overlap=10,tilesize=(30,30),k=5)  
plt.imshow(img2_stitch,cmap=plt.cm.gray)  
plt.show()
```


Original Image:



library has ntiles = 38115 each with npix = 900
row:

